

---

# Uncertainty-Driven Adaptive Frame Selection for Video Question Answering via Reinforcement Learning

Ishan Katpally, Om Mahesh · University of California, Santa Barbara · [ishankatpally@ucsb.edu](mailto:ishankatpally@ucsb.edu), [ommahesh@ucsb.edu](mailto:ommahesh@ucsb.edu)

---

## Abstract

Vision-language models (VLMs) waste compute on video by uniformly sampling frames regardless of the query. We propose a query-conditioned frame selection policy trained with PPO that treats video QA as active evidence-gathering: at each step, the policy keeps, skips, or stops watching based on accumulated evidence. The stopping criterion uses the VLM’s own output entropy, requiring no additional model. We evaluate on NExT-QA, EgoSchema, and ActivityNet-QA, where we expect adaptive stopping to yield question-dependent token budgets that outperform fixed-budget baselines at equal or lower average cost.

## 1 Introduction

Modern VLMs trained on image and video data, such as [Owen2.5-VL](#), process video by sampling frames and encoding them as visual tokens. A 32-frame clip can produce upward of 10,000 tokens, dominating the context window. Existing compression methods, like attention pruning ([FastV](#)), frame merging ([FrameFusion](#)), hierarchical compression ([LongVU](#)), compress based on visual redundancy before the question is considered, and always consume a fixed token budget regardless of question difficulty. We argue that video QA is fundamentally an active evidence-gathering problem: the model should keep watching when it is uncertain, and stop as soon as it has gathered sufficient temporal evidence to answer confidently. This exploration-exploitation framing of exploring more frames vs. exploiting current evidence yields question-adaptive budgets that are both more efficient and more principled than fixed compression schemes.

## 2 Related Work

**Static video token compression.** Attention pruning ([FastV](#)), similarity merging ([FrameFusion](#)), and hierarchical compression ([LongVU](#); [VideoChat-Flash](#)) all compress based on visual redundancy alone and apply a fixed budget regardless of query difficulty. Our method conditions both frame selection and stopping on the input query and model uncertainty.

**RL for visual efficiency in VLMs.** [VisionThink](#) uses RL to request higher-resolution images when needed. [Quickviewer](#) applies RL to video cube compression. Both use fixed token budgets and neither models uncertainty over the temporal evidence accumulated so far. Our key distinction is an explicit exploration-exploitation tradeoff: the policy learns when it has seen enough, producing variable-length frame sequences driven by answer confidence rather than a predetermined budget.

## 3 Proposed Methodology

**Problem setup.** Given video  $V = \{f_1, \dots, f_{|V|}\}$  and query  $q$ , sequentially select frames to observe until the policy emits a stop action, then answer  $q$  using only the retained tokens, minimizing both errors and total frames consumed.

**MDP formulation.** At step  $t$  the state is the query embedding concatenated with a mean-pool of accepted frame embeddings and the current VLM output entropy  $H_t$  (computed from a partial forward pass over accepted frames so far). The action space is ternary: keep, skip, or stop. Upon stopping, the frozen VLM generates a final answer. The reward is  $R = \mathbb{1}[\text{correct}] - \lambda \cdot |S|/N$ , with  $\lambda \in \{0.05, 0.1, 0.2\}$ . Including  $H_t$  in the state directly incentivizes the policy to stop when the VLM is already confident and to keep exploring when it is not.

**Selector architecture.** Frame embeddings and query encodings come from the frozen [Owen2.5-VL-7B](#) encoders. A 2-layer transformer selector takes the current state and outputs keep/skip/stop logits. We warm-start via supervised imitation of FastV scores before PPO fine-tuning with [HuggingFace TRL](#). Entropy  $H_t$  is computed with no gradient through the frozen VLM at each step.

## 4 Evaluation

**Benchmarks.** [NExT-QA](#) (causal and temporal reasoning), [EgoSchema](#) (long-form egocentric, ~3 min), and [ActivityNet-QA](#) (open-ended activity understanding).

**Metrics.** Top-1 accuracy (primary); average token retention rate  $|S|/N$ ; accuracy-at-budget curves; and distribution of stopping points per question type, showing that easy questions trigger early stopping and hard ones do not.

**Baselines.** Uniform sampling at 8, 16, 32 frames; [FastV](#) pruning; [FrameFusion](#) merging; [VisionThink](#)-style fixed-budget RL; full-frame uncompressed VLM as upper bound.

**Interpretability analysis.** We visualize stopping-point distributions across NExT-QA question types (causal, temporal, descriptive) and plot entropy  $H_t$  over time per category, testing whether the policy learns that causal questions require more temporal exploration than descriptive ones.